

```
## Objective
#Our goal is to predict the price of a BMW car considering their unique characteristics

#We will import google.colab as well as some libraries

from google.colab import drive
drive.mount('/content/drive')

import pandas as pd
import numpy as np
from datetime import datetime
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder,MinMaxScaler
from sklearn import model_selection # model assesment and model selection strategies
from sklearn import metrics # model evaluation metrics

import folium
import plotly.express as px

📁 Mounted at /content/drive

#Reading the .csv and take a first look into the dataset
df_bmw=pd.read_csv("/content/drive/MyDrive/ADB/Master Data Science/2. Archivos de clases/Proyecto #1 - BMW/bmw_pricing_v2.csv")
df_bmw.head()
```

	marca	modelo	km	potencia	fecha_registro	tipo_gasolina	color	tipo_coche
0	BMW	118	140411.0	100.0	2012-02-01	diesel	black	convertible
1	BMW	M4	13929.0	317.0	2016-04-01	petrol	grey	convertible
2	BMW	320	183297.0	120.0	2012-04-01	diesel	white	convertible
3	BMW	420	128035.0	135.0	2014-07-01	diesel	red	convertible
4	BMW	425	97097.0	160.0	2014-12-01	diesel	silver	convertible

```
df_bmw.info()
```

```
# We'll delete the column named "marca" due to there is any kind of information inside
# We'll round the column "potencia", change the type of "fecha"
# Out TARGET will be "precio"
# We'll change some variables from type object to boolean
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4843 entries, 0 to 4842
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   marca                                4841 non-null   object
1   modelo                              4840 non-null   object
2   km                                   4841 non-null   float64
3   potencia                             4842 non-null   float64
4   fecha_registro                       4842 non-null   object
5   tipo_gasolina                        4838 non-null   object
6   color                                4831 non-null   object
7   tipo_coche                           4834 non-null   object
8   volante_regulable                   4839 non-null   object
9   aire_acondicionado                  4841 non-null   object
10  camara_trasera                       4841 non-null   object
11  asientos_traseros_plegables          4839 non-null   object
12  elevalunas_electrico                 4841 non-null   object
13  bluetooth                            4839 non-null   object
14  gps                                   4843 non-null   bool
15  alerta_lim_velocidad                 4841 non-null   object
16  precio                               4837 non-null   float64
17  fecha_venta                          4842 non-null   object
dtypes: bool(1), float64(3), object(14)
memory usage: 648.1+ KB
```

```
df_bmw.describe().T #Another first look into the dataset
```

```
#There are some "km" in negative, as well as with "1M"
#There are some "potential" with zero
#There are some prices really cheaper
```

	count	mean	std	min	25%	50%
km	4841.0	140959.347862	60208.534313	-64.0	102884.0	141080.0

df\_bmw.isnull().sum() #There are some nulls that we'll delete from the dataset

marca	2
modelo	3
km	2
potencia	1
fecha_registro	1
tipo_gasolina	5
color	12
tipo_coche	9
volante_regulable	4
aire_acondicionado	2
camara_trasera	2
asientos_traseros_plegables	4
elevalunas_electrico	2
bluetooth	4
gps	0
alerta_lim_velocidad	2
precio	6
fecha_venta	1
dtype: int64	


len(df\_bmw[df\_bmw.duplicated()]) #We'll identify if there is some duplicated values

0

#### ANALYSIS EACH COLUMN ####

## PRICE ##

df\_bmw['precio'].hist()



sns.boxplot(x=df\_bmw["precio"])

```
<Axes: xlabel='precio'>

len(df_bmw[df_bmw['precio'] < 500]) #There are 15 cars with value below from 500€

15

df_bmw[df_bmw['precio'] < 500]

#Models: 320, 318, 520, 116, 525, 316, X3
```

	marca	modelo	km	potencia	fecha_registro	tipo_gasolir
565	BMW	320	179358.0	120.0	2013-06-01	dies
630	BMW	318	147558.0	105.0	2014-11-01	Na
879	BMW	318	134156.0	105.0	2014-06-01	dies
1255	BMW	320	170381.0	135.0	2013-07-01	dies
1513	BMW	520	358332.0	100.0	2000-10-01	dies
1558	BMW	520	358333.0	100.0	2000-10-01	dies
1832	BMW	116	174524.0	85.0	2014-07-01	dies
2473	BMW	525	230578.0	85.0	1997-07-01	dies
2574	BMW	525	229880.0	85.0	1997-07-01	dies
2611	BMW	525	230264.0	85.0	1997-07-01	dies
2829	BMW	525	439060.0	105.0	1996-10-01	dies
3062	BMW	318	98097.0	85.0	1994-01-01	petr


```
df_bmw.drop(df_bmw[df_bmw['precio'] <500].index, inplace=True) #Eliminating cars with values < 500

len(df_bmw[df_bmw['precio'] > 100000]) #There are 4 cars with value > 100.000€

2

df_bmw.drop(df_bmw[df_bmw['precio'] > 100000].index, inplace=True) #Eliminating cars with values > 100.000

df_bmw['precio'].hist() #Validamos que aún tiene una distribución a la derecha, pero luego normalizaremos el TARGET
```



## MARCA ##

```
df_bmw['marca'].value_counts() #There is only one data
```

```
df_bmw['marca'].value_counts() #there is only one data
```

```
BMW      4824
Name: marca, dtype: int64
```

```
del(df_bmw['marca']) #Eliminating the column
```

```
## Tipo coche ##
```

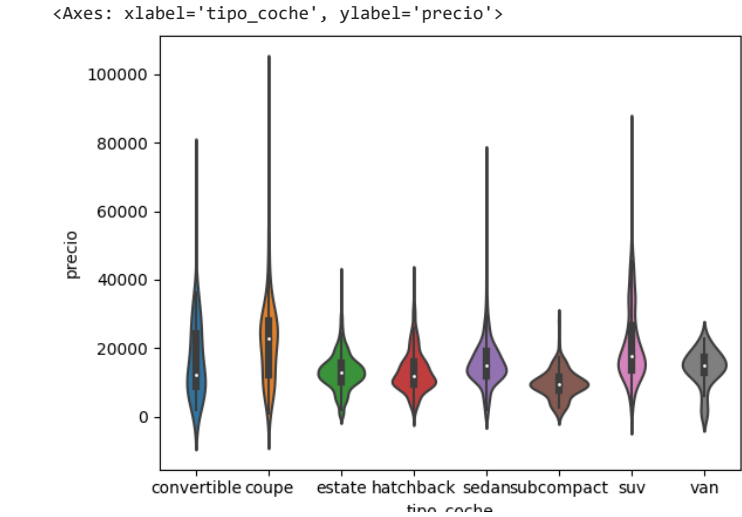
```
df_bmw['tipo_coche'].value_counts() #Taking a deep looking
```

```
#First we'll identify the relation between price and type of car
```

```
estate      1598
sedan       1160
suv         1054
hatchback   698
subcompact  113
coupe       104
convertible  47
van         43
Name: tipo_coche, dtype: int64
```

```
sns.violinplot(x="tipo_coche", y="precio", data=df_bmw) #NO NUMERICAL VARIABLE
```

```
#We'll hold on the analysis of this column
```



```
## Modelo ##
```

```
df_bmw[df_bmw['modelo'].isnull()]
```

	modelo	km	potencia	fecha_registro	tipo_gasolina	color
173	NaN	146338.0	105.0	2014-01-01	diesel	blanco
4766	NaN	115566.0	105.0	2014-01-01	diesel	silver

```
len(df_bmw["modelo"].value_counts())
```

```
76
```

```
df_bmw["modelo"].value_counts()
```

```
320      750
520      631
318      565
X3       436
116      357
...
M135      1
225      1
i8        1
630      1
```

```
214 Gran Tourer      1
Name: modelo, Length: 76, dtype: int64
```

```
df_bmw[(df_bmw["modelo"] == "i8") & (df_bmw["precio"]>50000)]
```

```
#Model i8 & price > 50.000 --> outlier
```

modelo	km	potencia	fecha_registro	tipo_gasolina	color
[Redacted]					

```
df_bmw = df_bmw.drop(df_bmw[(df_bmw["modelo"] == "i8") & (df_bmw["precio"] > 50000)].index)
```

```
df_bmw[(df_bmw["modelo"] == "X3") & (df_bmw["precio"]>50000)]
#Model X3 & price > 50.000 --> outlier
```

modelo	km	potencia	fecha_registro	tipo_gasolina	color	tipo_gasolina
[Redacted]						

```
df_bmw = df_bmw.drop(df_bmw[(df_bmw["modelo"] == "X3") & (df_bmw["precio"] > 50000)].index)
```

```
df_bmw[(df_bmw["modelo"] == "i8") & (df_bmw["precio"]>1000)]
#Model i8 & price > 1.000 --> unique
```

modelo	km	potencia	fecha_registro	tipo_gasolina	color	tipo_gasolina
[Redacted]						

```
df_bmw = df_bmw.drop(df_bmw[(df_bmw["modelo"] == "i8") & (df_bmw["precio"] > 1000)].index)
```

```
## Tipo gasolina ##
```

```
df_bmw[df_bmw['tipo_gasolina'].isnull()]
```

	modelo	km	potencia	fecha_registro	tipo_gasolina	color
<b>82</b>	420	54993.0	135.0	2014-03-01	NaN	black
<b>185</b>	320	186697.0	135.0	2012-11-01	NaN	white
<b>444</b>	318	111622.0	100.0	2013-01-01	NaN	black
[Redacted]						

```
df_bmw["tipo_gasolina"].value_counts()
```

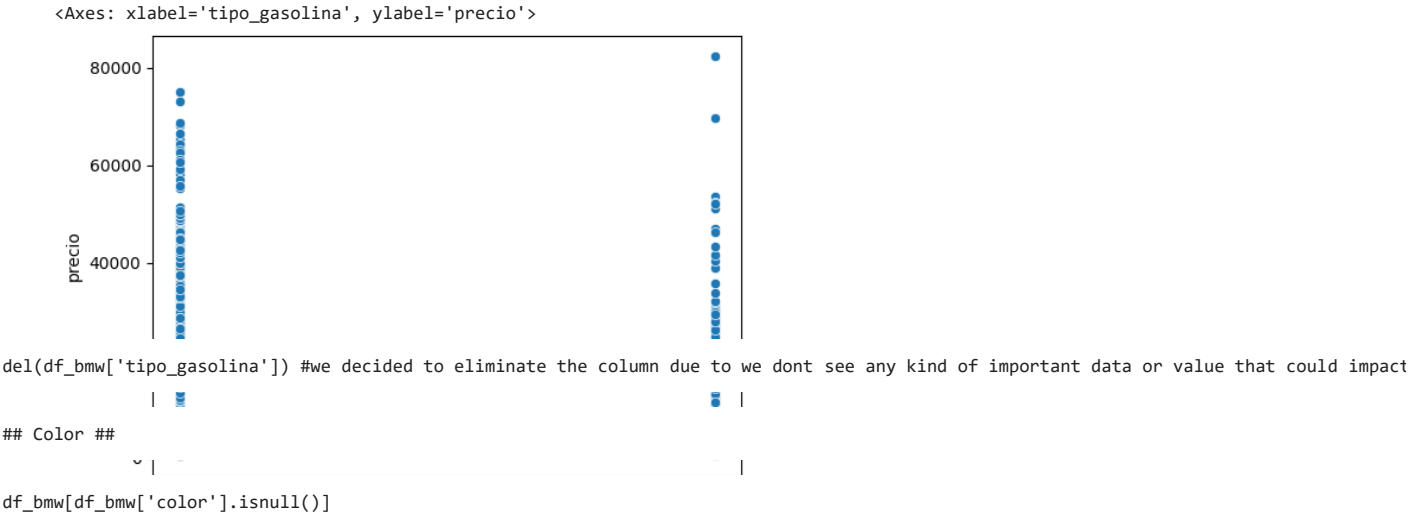
```
diesel      4618
petrol      188
hybrid_petrol    7
Diesel       5
electro       3
Name: tipo_gasolina, dtype: int64
```

```
df_bmw["tipo_gasolina"]=df_bmw["tipo_gasolina"].replace({'diesel':'Diesel'}) #group
df_bmw["tipo_gasolina"] = df_bmw["tipo_gasolina"].replace(['Hybrid Petrol', 'Electro'], 'petrol') #group
df_bmw["tipo_gasolina"] = df_bmw["tipo_gasolina"].replace(['hybrid_petrol', 'electro'], 'petrol') #group
```

```
df_bmw["tipo_gasolina"].value_counts()
```

```
Diesel      4623
petrol      198
Name: tipo_gasolina, dtype: int64
```

```
sns.scatterplot(x="tipo_gasolina", y="precio", data=df_bmw) #NUMERIC VARIABLE
```



	modelo	km	potencia	fecha_registro	color	tipo_coche
239	318	132731.0	100.0	2013-09-01	NaN	estate
834	318	148429.0	100.0	2013-06-01	NaN	estate
855	318	139736.0	100.0	2009-02-01	NaN	estate
864	318	157661.0	100.0	2013-06-01	NaN	estate
884	320	145981.0	122.0	2013-07-01	NaN	estate
904	320	126425.0	120.0	2013-07-01	NaN	estate
939	520	153102.0	140.0	2015-04-01	NaN	estate
1569	318	191804.0	100.0	2013-10-01	NaN	estate
1591	320	130624.0	120.0	2013-07-01	NaN	estate

```
df_bmw['color'].value_counts()
```

```
black      1627
grey       1166
blue        702
white       536
brown       341
silver      325
red         51
beige       41
green       18
orange       6
Name: color, dtype: int64
```

```
sns.scatterplot(x="color", y="precio", data=df_bmw)
```

```

<Axes: xlabel='color', ylabel='precio'>
del(df_bmw['color']) #we decided to eliminate the column due to we dont see any kind of important data or value that could impact on the

### COPY ###

df_bmw1 = df_bmw.copy() #we generate a copy so we can eliminate the nulls
df_bmw1.dropna(inplace = True)

df_bmw1.isnull().sum() #Checking if there are some nulls after the dropping

modelo      0
km           0
potencia    0
fecha_registro  0
tipo_coche  0
volante_regulable  0
aire_acondicionado  0
camara_trasera  0
asientos_traseros_plegables  0
elevalunas_electrico  0
bluetooth   0
gps         0
alerta_lim_velocidad  0
precio      0
fecha_venta 0
dtype: int64

### Fecha registro ###

# Change the type of the column to a date format
df_bmw1['fecha_registro']=pd.to_datetime(df_bmw1['fecha_registro'])

# Take the year from the column
df_bmw1['anio_registro']=df_bmw1['fecha_registro'].apply(lambda x:datetime.strftime(x,'%Y'))
df_bmw1['mes_registro'] = df_bmw1['fecha_registro'].apply(lambda x: datetime.strftime(x, '%m'))

df_bmw1["anio_registro"].value_counts() #Checking the format

2013    1527
2014    1270
2012     817
2015     310
2011     214
2010     103
2008     102
2009      85
2016      82
2006      67
2007      58
2005      52
2004      26
2001      17
2003      16
2017      11
2002       8
2000       8
1999       3
1998       2
1996       1
1995       1
1990       1
1997       1
Name: anio_registro, dtype: int64

df_bmw1["mes_registro"].value_counts()

03      464
01      452
10      446
07      445
06      442
05      435
09      416
04      406
02      393
11      343
08      331

```

```

12      209
Name: mes_registro, dtype: int64

# Change the type of the column to a date format
df_bmw1['fecha_venta']=pd.to_datetime(df_bmw1['fecha_venta'])

# Take the year from the column
df_bmw1['anio_venta']=df_bmw1['fecha_venta'].apply(lambda x:datetime.strptime(x,'%Y'))
df_bmw1['mes_venta'] = df_bmw1['fecha_venta'].apply(lambda x: datetime.strptime(x, '%m'))

df_bmw1["anio_venta"].value_counts()

2018      4778
2007         1
2010         1
2009         1
2008         1
Name: anio_venta, dtype: int64

del(df_bmw1['anio_venta']) #We'll eliminate the year of the sale due to its irrelevant

df_bmw1["mes_venta"].value_counts()

05      804
03      726
04      688
06      599
07      532
08      519
02      490
09      220
01      204
Name: mes_venta, dtype: int64

del(df_bmw1['fecha_registro']) #Eliminating the column

del(df_bmw1['fecha_venta'])

### Booleans type ###

#We'll change the type of the following variables

variables_booleanas = ["tipo_coche", "volante_regulable", "aire_acondicionado", "camara_trasera", "asientos_traseros_plegables", "elevacur

for i in range(len(variables_booleanas)):
    variables_booleanas[i] = bool(variables_booleanas[i])

### Volante_regulable ###

df_bmw1["volante_regulable"].value_counts() #We consider it a great variable, with a good distribution

True      2638
False     2144
Name: volante_regulable, dtype: int64

sns.violinplot(x="volante_regulable", y="precio", data=df_bmw1) #NO NUMERIC

```



<Axes: xlabel='volante\_regulable', ylabel='precio'>

80000


### Aire\_acondicionado ###

df\_bmw1["aire\_acondicionado"].value\_counts() #80% - 20%

```
True      3807
False     975
Name: aire_acondicionado, dtype: int64
```

sns.violinplot(x="aire\_acondicionado", y="precio", data=df\_bmw1) #NO NUMERIC

<Axes: xlabel='aire\_acondicionado', ylabel='precio'>




### Camara trasera ###

df\_bmw1["camara\_trasera"].value\_counts() #80% - 20%

```
False     3817
True       965
Name: camara_trasera, dtype: int64
```

sns.violinplot(x="camara\_trasera", y="precio", data=df\_bmw1) #NO NUMERIC

<Axes: xlabel='camara\_trasera', ylabel='precio'>



### Asientos\_traseros\_plegables ###

```
df_bmw1["asientos_traseros_plegables"].value_counts() #80-20
```

```
False    3824
True      958
Name: asientos_traseros_plegables, dtype: int64
```

```
### Elevalunas_electrico ###
```

```
df_bmw1["elevalunas_electrico"].value_counts() #50-50
```

```
False    2575
True     2207
Name: elevalunas_electrico, dtype: int64
```

```
### Bluetooth ###
```

```
df_bmw1["bluetooth"].value_counts() #70-30
```

```
False    3622
True     1160
Name: bluetooth, dtype: int64
```

```
### GPS ###
```

```
df_bmw1["gps"].value_counts() #90-10
```

```
True     4462
False     320
Name: gps, dtype: int64
```

```
### Alerta_lim_velocidad ###
```


```
df_bmw1["alerta_lim_velocidad"].value_counts() #40-60
```

```
True     2588
False    2194
Name: alerta_lim_velocidad, dtype: int64
```

```
### KM ###
```

```
sns.scatterplot(x="km", y="precio", data=df_bmw1) #NUMERIC
```


```
<Axes: xlabel='km', ylabel='precio'>
```



```
df_bmw1[df_bmw1["km"]<100] #Looking for undervalued data
```

```
#We'll eliminate the negative rows
```

```
modelo    km  potencia  tipo_coche  volante_regulable  aire_ac
```



```
df_bmw1.drop(df_bmw1[df_bmw1['km'] < 0].index, inplace=True)
```

```
df_bmw1[df_bmw1["km"]>400000] #Looking for overvalued data
```

```
#We'll eliminate KM >1.000.000
```


	modelo	km	potencia	tipo_coche	volante_regulable	air
557	520	484615.0	120.0	estate	True	
1573	320	400654.0	110.0	estate	False	
2350	318	477571.0	85.0	hatchback	False	
3198	320	405816.0	100.0	sedan	False	

```
df_bmw1.drop(df_bmw1[df_bmw1['km'] > 1000000].index, inplace=True)
```

```
sns.scatterplot(x="km", y="precio", data=df_bmw1) #NUMERIC
```

```
#Here we have a correlation: - KM + price
```

```
<Axes: xlabel='km', ylabel='precio'>
```



```
### POWER ###
```

```
df_bmw2 = df_bmw1.copy() #Generating a copy
```

```
sns.scatterplot(x="potencia", y="precio", data=df_bmw2) #NUMERIC
```

```
#There could be a correlation: + price + power
```

```
<Axes: xlabel='potencia', ylabel='precio'>

df_bmw2[df_bmw2["potencia"]<70] #There are 5 rows with power < 70




|      | modelo | km       | potencia | tipo_coche | volante_regulable | aire |
|------|--------|----------|----------|------------|-------------------|------|
| 1796 | i3     | 152328.0 | 25.0     | hatchback  | False             |      |
| 1925 | i3     | 152470.0 | 25.0     | hatchback  | False             |      |
| 2390 | 318    | 170529.0 | 66.0     | hatchback  | False             |      |
| 2771 | 316    | 146951.0 | 66.0     | sedan      | False             |      |



df_bmw2.drop(df_bmw2[df_bmw2["potencia"] < 70].index, inplace=True)

sns.scatterplot(x="potencia", y="precio", data=df_bmw2)

<Axes: xlabel='potencia', ylabel='precio'>
```



```
df_bmw2[df_bmw2["potencia"]>350] #Looking for outliers



|      | modelo | km       | potencia | tipo_coche | volante_regulable | aire |
|------|--------|----------|----------|------------|-------------------|------|
| 3601 | M5     | 150187.0 | 412.0    | sedan      | True              |      |


```

```
df_bmw2[df_bmw2["modelo"] == "X6 M"] #We'll eliminate the row with power = 423
#Possible outlier if we compare it with the rest of the rows with same model
```

	modelo	km	potencia	tipo_coche	volante_regulable	aire
3829	X6 M	39725.0	280.0	suv	False	
3986	X6 M	115569.0	280.0	suv	True	
4109	X6 M	67798.0	190.0	suv	True	
4146	X6 M	2970.0	423.0	suv	True	
4166	X6 M	53221.0	180.0	suv	True	
4282	X6 M	90157.0	190.0	suv	True	

```
df_bmw2.drop(df_bmw2[df_bmw2["potencia"] > 420].index, inplace=True)

df_bmw2[df_bmw2["potencia"]>350] #Checking
```

	modelo	km	potencia	tipo_coche	volante_regulable	aire


```
### TARGET ###
```

```
df_bmw3 = df_bmw2.copy()
```

```
sns.boxplot(x=df_bmw3["precio"])
```


```
#Looking for outliers
```

```
<Axes: xlabel='precio'>
```



```
df_bmw3["precio"].hist()
```


```
<Axes: >
```



```
df_bmw3['precioLN'] = df_bmw3['precio'].apply(lambda x: np.log1p(x))
TARGET_LN = 'precioLN'
```

```
df_bmw3["precioLN"].hist()
```

&lt;Axes: &gt;



```
del(df_bmw3["precioLN"])
```

```
|| | | | |
```

```
### TRANSFORMATIONS: OHE & MINMAXSCALER ###
```

```
|| | | | |
```

```
df_bmw4 = df_bmw3.copy()
```

```
|| | | | |
```

```
df_bmw4.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 4774 entries, 0 to 4841
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	modelo	4774 non-null	object
1	km	4774 non-null	float64
2	potencia	4774 non-null	float64
3	tipo_coche	4774 non-null	object
4	volante_regulable	4774 non-null	object
5	aire_acondicionado	4774 non-null	object
6	camara_trasera	4774 non-null	object
7	asientos_traseros_plegables	4774 non-null	object
8	elevallunas_electrico	4774 non-null	object
9	bluetooth	4774 non-null	object
10	gps	4774 non-null	bool
11	alerta_lim_velocidad	4774 non-null	object
12	precio	4774 non-null	float64
13	anio_registro	4774 non-null	object
14	mes_registro	4774 non-null	object
15	mes_venta	4774 non-null	object

```
dtypes: bool(1), float64(3), object(12)
```

```
memory usage: 601.4+ KB
```

```
df_bmw4["potencia"] = MinMaxScaler().fit_transform(df_bmw4["potencia"].values.reshape(-1,1)) #We'll replace with MMS the power column
```

```
df_bmw4["km"] = MinMaxScaler().fit_transform(df_bmw4["km"].values.reshape(-1,1)) #Same action with KM
```

```
# We create a boolean list
```

```
variables_booleanas = ["tipo_coche", "volante_regulable", "aire_acondicionado", "camara_trasera", "asientos_traseros_plegables", "elevallur
```

```
for column in variables_booleanas:
```

```
    # OHE
```

```
    encoded_columns = pd.get_dummies(df_bmw4[column], prefix=column)
```

```
    # Concat the DataFrame with the transformed column
```

```
    df_bmw4 = pd.concat([df_bmw4, encoded_columns], axis=1)
```

```
# Eliminating the original boolean columns
```

```
df_bmw4 = df_bmw4.drop(variables_booleanas, axis=1)
```

```
df_bmw4.head()
```

	km	potencia	precio	tipo_coche_convertible	tipo_coche_co
0	0.289039	0.087719	11300.0		1
1	0.027787	0.722222	69700.0		1
2	0.377621	0.146199	10200.0		1
3	0.263476	0.190058	25100.0		1
4	0.199573	0.263158	33400.0		1

```
### MODELING ###
```

```
#We decide to use aleatory to split the model
```

```
p_dev = 0.70 # % of train
```

```
df_bmw4['is_train'] = np.random.uniform(0, 1, len(df_bmw4)) <= p_dev
dev_df_bmw4, val_df_bmw4 = df_bmw4[df_bmw4['is_train']==True], df_bmw4[df_bmw4['is_train']==False]
df_bmw4 = df_bmw4.drop('is_train', 1)
```

```
print("Ejemplos usados para entrenar: ", len(dev_df_bmw4))
print("Ejemplos usados para validación: ", len(val_df_bmw4))
```

```
Ejemplos usados para entrenar: 3277
```

```
Ejemplos usados para validación: 1497
```

```
<ipython-input-117-55dee41ae9a5>:7: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the arg
df_bmw4 = df_bmw4.drop('is_train', 1)
```

```
### Asignación de atributos y target a las variables X e Y ###
```

```
dev_df_bmw4_X = dev_df_bmw4.drop('precio', axis=1)
dev_df_bmw4_y = dev_df_bmw4[['precio']]
```

```
val_df_bmw4_X = val_df_bmw4.drop('precio', axis=1)
val_df_bmw4_y = val_df_bmw4[['precio']]
```

```
### Random hold out ###
```

```
X_train, X_test, y_train, y_test = model_selection.train_test_split(
    dev_df_bmw4_X, # X
    dev_df_bmw4_y, # y
    test_size = 0.20,
    random_state = 1279
)
```

```
### Previous checks ###
```

```
dev_df_bmw4_X.head().T
```

	1	3	4	5	7
km	0.027787	0.263476	0.199573	0.313703	0.237709
potencia	0.722222	0.190058	0.263158	0.453216	0.102338
tipo_coche_convertible	1	1	1	1	1
tipo_coche_coupe	0	0	0	0	0
tipo_coche_estate	0	0	0	0	0
...	...	...	...	...	...
modelo_X6	0	0	0	0	0
modelo_X6 M	0	0	0	0	0
modelo_Z4	0	0	0	0	0

```
X_train.info(verbose=False)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2621 entries, 3207 to 1464
Columns: 146 entries, km to is_train
dtypes: bool(1), float64(2), uint8(143)
memory usage: 430.0 KB
```

```
X_test.info(verbose=False)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 656 entries, 1347 to 846
Columns: 146 entries, km to is_train
dtypes: bool(1), float64(2), uint8(143)
memory usage: 107.6 KB
```

```
X_train.describe().T.head()
```

	count	mean	std	min	25%
km	2621.0	0.287988	0.120934	0.000281	0.208839
potencia	2621.0	0.171648	0.110721	0.000000	0.087719
tipo_coche_convertible	2621.0	0.011446	0.106392	0.000000	0.000000
tipo_coche_coupe	2621.0	0.020221	0.140783	0.000000	0.000000

```
X_test.describe().T.head()
```

	count	mean	std	min	25%
km	656.0	0.290215	0.117748	0.000475	0.220167
potencia	656.0	0.173602	0.123289	0.000000	0.087719
tipo_coche_convertible	656.0	0.006098	0.077908	0.000000	0.000000
tipo_coche_coupe	656.0	0.030488	0.172056	0.000000	0.000000

```
y_train.describe().T.head()
```

	count	mean	std	min	25%	50%

```
y_test.describe().T.head()
```

	count	mean	std	min	25%	50%

```
### MODEL DEFINITION - REGRESSION ###
```

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import export_graphviz
import graphviz
import pydotplus
```

```
!conda install python-graphviz -y
!conda install pydot -y
```

```
/bin/bash: line 1: conda: command not found
/bin/bash: line 1: conda: command not found
```

```
X_train, X_test, y_train, y_test = train_test_split(dev_df_bmw4_X, dev_df_bmw4_y, test_size=0.3, random_state=23)
```

```
lin_reg= LinearRegression()
lin_reg.fit(X_train,y_train)
```

```
LinearRegression()
LinearRegression()
```

```
y_pred = lin_reg.predict(X_test)
```

```
y_test.head()
```



```

    precio
4747  8000.0
-----
df_resultados = pd.DataFrame({'Actual': y_test, 'Predicted':y_pred})

-----
ValueError                                Traceback (most recent
call last)
<ipython-input-137-fa470504ef3a> in <cell line: 1>()
----> 1 df_resultados = pd.DataFrame({'Actual': y_test,
'Predicted':y_pred})

-----
3 frames
/usr/local/lib/python3.10/dist-
packages/pandas/core/internals/construction.py in
_extract_index(data)
    651         raw_lengths.append(len(val))
    652         elif isinstance(val, np.ndarray) and val.ndim >
1:

```

```
print ("MAE", metrics.mean_absolute_error(y_test, y_pred))
```

```
MAE 16745316344196.914
```

```
print ("MSE", metrics.mean_squared_error(y_test, y_pred))
```

```
MSE 3.917628648694983e+28
```

```
print("RMSE",np.sqrt(metrics.mean_squared_error(y_test, y_pred, squared = False)))
```

```
RMSE 14068759.860373845
```

```
df_resultados["dif"]=df_resultados["Predicted"]-df_resultados["Actual"]
```

```

-----
NameError                                Traceback (most recent
call last)
<ipython-input-141-5d3445ad96f7> in <cell line: 1>()
----> 1 df_resultados["dif"]=df_resultados["Predicted"]-
df_resultados["Actual"]

```

```
df_resultados.sort_values(by="dif")
```

```

-----
NameError                                Traceback (most recent
call last)
<ipython-input-142-e57d28c6b154> in <cell line: 1>()
----> 1 df_resultados.sort_values(by="dif")

NameError: name 'df_resultados' is not defined

```

```
df_resultados.hist("dif")
```

```

-----
NameError                                Traceback (most recent
call last)
<ipython-input-143-1bbb1a882813> in <cell line: 1>()
----> 1 df_resultados.hist("dif")

NameError: name 'df_resultados' is not defined

```

```
### Model definition - RandomForestRegression ###
```

```
from sklearn.ensemble import RandomForestRegressor
```

```

modelo_rf = RandomForestRegressor(n_estimators = 20, random_state = 2022)
modelo_rf.fit(dev_df_bmw4_X, dev_df_bmw4_y)

```

```
<ipython-input-146-43a17aca18df>:2: DataConversionWarning: A column-
modelo_rf.fit(dev_df_bmw4_X, dev_df_bmw4_y)
```

```
RandomForestRegressor
```

```
predicciones_rf = modelo_rf.predict(X_test)
predicciones_rf
```

```
26085., 19245., 13975., 10120., 39870., 5240., 13320., 8875.,
15915., 9920., 19860., 9665., 39855., 13900., 4455., 14710.,
12715., 10385., 10480., 6855., 26345., 16800., 50715., 10365.,
7655., 13470., 14080., 14110., 19210., 13130., 19530., 13600.,
11110., 18145., 14840., 17675., 12420., 16540., 22160., 10035.,
13400., 29695., 19515., 10310., 36895., 15140., 26260., 19290.,
18890., 11695., 12815., 11110., 37820., 10795., 3710., 13790.,
14215., 19820., 9905., 10760., 16990., 6435., 14345., 22205.,
14150., 12805., 14830., 20645., 11545., 10010., 3335., 6455.,
7655., 19730., 19160., 19480., 5165., 15625., 16565., 15335.,
8670., 25415., 12975., 22820., 13245., 15620., 17375., 13330.,
13350., 11220., 14340., 8510., 10705., 15690., 14595., 27955.,
12800., 12450., 12945., 11825., 29825., 13965., 12395., 28330.,
14090., 13135., 13840., 10110., 21360., 21200., 13085., 17295.,
14395., 19770., 18580., 9120., 23505., 13235., 14495., 14705.,
12475., 19200., 13760., 12760., 20850., 10155., 8735., 7070.,
4615., 1915., 12370., 7790., 15185., 12985., 12945., 31170.,
29970., 12715., 17100., 37575., 9110., 15390., 9595., 13775.,
11170., 9725., 14405., 13130., 12930., 12165., 17385., 24815.,
20805., 8865., 12715., 26480., 20145., 12865., 26555., 14880.,
13810., 1335., 20205., 23720., 11330., 16375., 14330., 18360.,
18495., 19575., 6325., 7880., 10515., 24380., 18380., 22600.,
22090., 14530., 11525., 15340., 11825., 14985., 15925., 31365.,
15355., 13645., 19425., 20235., 11635., 11705., 17740., 30445.,
12650., 15540., 11065., 10240., 1300., 17960., 12355., 8320.,
12160., 8670., 20490., 27325., 11680., 13925., 19875., 10710.,
19950., 25175., 14215., 9835., 12775., 18375., 18180., 17205.,
22360., 7035., 11310., 4320., 11685., 29230., 17005., 15445.,
12680., 12960., 20155., 38475., 18790., 9850., 10685., 15255.,
19290., 20675., 10295., 17825., 18250., 15135., 15130., 33940.,
12985., 13570., 10930., 36620., 14590., 17690., 20065., 30270.,
47785., 25890., 15260., 18135., 5415., 11805., 11265., 9080.,
5055., 10775., 12095., 24760., 14320., 22610., 30135., 17190.,
28440., 1290., 26175., 15705., 22170., 13235., 10020., 11040.,
20260., 9505., 14375., 10385., 28690., 11320., 15950., 13675.,
6690., 11305., 18085., 15530., 11920., 5335., 28005., 9245.,
23395., 12590., 11380., 22165., 14405., 13105., 13720., 11965.,
10555., 15540., 6030., 9215., 8545., 20955., 38705., 29915.,
14940., 9185., 17160., 10660., 23890., 14625., 14720., 8890.,
21100., 29570., 28585., 11905., 33115., 7830., 17820., 9630.,
25705., 14170., 18400., 11780., 28270., 20885., 12515., 11330.,
14240., 16855., 16585., 9940., 22375., 10950., 11735., 18595.,
9965., 13075., 11725., 8525., 12035., 13635., 9250., 24090.,
22160., 22310., 17460., 13375., 12070., 18000., 11085., 14495.,
18150., 16960., 16000., 7720., 9375., 18185., 17460., 18340.,
9950., 6865., 4225., 18040., 16525., 21835., 22105., 13225.,
16595., 12155., 13105., 19630., 21740., 16100., 12615., 5610.,
14865., 16540., 6200., 19360., 14665., 11880., 20855., 14445.,
19045., 21085., 13320., 13040., 18555., 14195., 11915., 39200.,
11300., 4635., 12925., 17595., 16875., 19670., 26295., 12405.,
8350., 11975., 9365., 14030., 9430., 11685., 12460., 16220.,
11595., 18805., 13210., 11825., 3300., 7690., 12475., 13795.,
19815., 33670., 10980., 10450., 13140., 41315., 11665., 21435.,
12875., 14360., 11975., 11640., 14885., 9535., 9335., 14060.,
12405., 16480., 15720., 14965., 10275., 5380., 10230., 35120.,
10120., 20100., 17010., 54620., 17005., 12470., 25780., 9415.,
7775., 14365., 10790., 13260., 14300., 13895., 13250., 9990.,
29705., 25690., 17995., 11020., 22305., 13010., 12340., 21065.,
13455., 11280., 5250., 9935., 26680., 12560., 25575., 12405..
```

```
comparaciones = pd.DataFrame(X_test)
comparaciones = comparaciones.assign(Precio_Real = y_test)
comparaciones = comparaciones.assign(Precio_Prediccion = predicciones_rf.flatten().tolist())
print(comparaciones)
```

```
...      ...      ...      ...      ...
4691  0.188138  0.248538      0      0
3007  0.373422  0.190058      0      0
1496  0.238442  0.146199      0      0
2245  0.188894  0.087719      0      0
3590  0.371552  0.321637      0      0

      tipo_coche_estate  tipo_coche_hatchback  tipo_coche_sedan  \
1717                  1                    0                  0
2594                  0                    0                  1
```

2245	0	1	0		
3590	0	0	1		
	tipo_coche_subcompact	tipo_coche_suv	tipo_coche_van	...	modelo_X5 \
1717	0	0	0	...	0
2594	0	0	0	...	0
1880	0	0	0	...	0
2027	0	0	0	...	0
3293	0	0	0	...	0
...	...	...	...	...	...
4691	0	1	0	...	1
3007	0	0	0	...	0
1496	0	0	0	...	0
2245	0	0	0	...	0
3590	0	0	0	...	0
	modelo_X5 M	modelo_X5 M50	modelo_X6	modelo_X6 M	modelo_Z4 \
1717	0	0	0	0	0
2594	0	0	0	0	0
1880	0	0	0	0	0
2027	0	0	0	0	0
3293	0	0	0	0	0
...	...	...	...	...	...
4691	0	0	0	0	0
3007	0	0	0	0	0
1496	0	0	0	0	0
2245	0	0	0	0	0
3590	0	0	0	0	0
	modelo_i3	is_train	Precio_Real	Precio_Prediccion	
1717	0	True	8900.0	9045.0	
2594	0	True	11300.0	11740.0	
1880	0	True	17300.0	18060.0	
2027	0	True	10200.0	10015.0	
3293	0	True	12300.0	12030.0	
...	...	...	...	...	
4691	0	True	23800.0	23625.0	
3007	0	True	15400.0	16060.0	
1496	0	True	12700.0	10865.0	
2245	0	True	12300.0	13065.0	
3590	0	True	18200.0	17970.0	

### RMSE modelo de ar ###

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
rmse_rf = mean_squared_error(
    y_true = y_test,
    y_pred = predicciones_rf,
    squared = False
)
print(f"El error (rmse) de test es: {rmse_rf}")
```

El error (rmse) de test es: 1335.4106562907753